

Jupyter Part

Let's code

Remove duplicates and write here how many duplicates were.

Answer

- There were a total of 10001 records
- There were a total of 9617 unique records
- There were a total of 384 duplicates

Add code for first cluster here

Code

```
# Read the CSV file
myRdd = sc.textFile("gs://kamal_at_ie/lab2Dataset.csv")
```

```
# Check the total no. of records and pprint that
total = myRdd.count()
print("Total:" + str(total))
```

```
# Create a new RDD containing only unique keys
uniqueRdd = myRdd.distinct()
```

```
# Print the total no. of unique keys
unique = uniqueRdd.count()
print("Unique:" + str(unique))
```

```
# Get the total no. of duplicates and print the same
dups = total - unique
print("Dups:" + str(dups))
```

Number of duplicates: 384

Create new RDD with ids whose length is bigger than 3.

Save RDD as text in gs bucket you created before. Choose a different file name.

Code (Continued from before):

```
# Take the original RDD that we had created earlier and create a new RDD that contains the
keys of length bigger than 3
idsgt3 = myRdd.filter(lambda line: len(line) > 3)
```

```
# Do a sanity check
idsgt3.take(10)
```

```
# Print the no. of strings in the new RDD
print("Count of IDs bigger than length 3:" + str(idsgt3.count()))
```

```
# Save this file in storage
idsgt3.saveAsTextFile("gs://kamal_at_ie/idsgt3.csv")
```

Delete cluster 1.

DataLab part

Read text file generated by cluster 1 and add it to rdd_main.

```
rdd_main = sc.textFile("gs://kamal_at_ie/ids3.csv")
```

Get number of partitions and reduce the partitions to 2

```
# get no. of partitions in the rdd
```

```
rdd_main.getNumPartitions()
```

```
2
```

```
# Reduce the partitions to 2; # although, we should do only if
```

```
#the existing no. of partitions is > 2
```

```
rddCoalesed = rdd_main.coalesce(2)
```

```
# again check if the no. of partitions are <=2
```

```
rdd_main.getNumPartitions()
```

```
2
```

Obtain the length of all the ids added together.

For example if ids are:

[alex, pepe, domin] then total would be 3 from alex plus 4 from pepe plus 5 from domin so 12

```
# get the length of all IDs put together
```

```
recLenRdd = rddCoalesed.map(lambda rec: len(rec))
```

```
totalLength = recLenRdd.reduce(lambda a, b: a + b)
```

```
print("Total length of all strings:" + str(totalLength))
```

```
Total length of all strings:91545
```

Sort alphabetically and display the first 10 ids (order from a to z)

```
#Sort the IDs alphabetically and print the first 10
```

```
sortedRdd = rddCoalesed.map(lambda x: (x, 1)).sortByKey().keys()
```

```
sortedRdd.take(10)
```

```
[u'02strich',
```

```
u'0m4r',
```

```
u'0x860111',
```

```
u'0x90',
```

```
u'0xen',
```

```
u'0xhacker',
```

```
u'10K35H 5H4KY4',
```

```
u'14256424',
```

```
u'150GritSandpaper',
```

```
u'1ifbyLAN2ifbyC']
```

Save the file to your bucket with the name sorted_rddids.txt

```
sortedRdd.saveAsTextFile("gs://kamal_at_ie/sorted_rddids.txt")
```

From rdd_main obtain the number of ids that share the first two characters.

For example if we have ids: aax , aat, aaron, bbt we would have aa,3 and

bb,1

```
rddWordCount = rdd_main.map(lambda x: (x[:2], 1)).reduceByKey(lambda a, b: a + b)
```

```
# print the first 10
```

```
rddWordCount.take(10)
```

```
[(u'gw', 1),  
(u'gu', 7),  
(u'gs', 4),  
(u'ge', 7),  
(u'gc', 1),  
(u'ga', 13),  
(u'go', 11),  
(u'gm', 4),  
(u'gk', 1),  
(u'm_', 1)]
```

Add ALL your code here

```
# read the file that contains IDs with length bigger than 3, from the storage bucket
```

```
rdd_main = sc.textFile("gs://kamal_at_ie/ids3.csv")
```

```
# Do some sanity check
```

```
print(rdd_main.count())
```

```
# get no. of partitions in the rdd
```

```
rdd_main.getNumPartitions()
```

```
# Reduce the partitions to 2; # although, we should do only if the existing no. of partitions  
is > 2
```

```
rddCoalesed = rdd_main.coalesce(2)
```

```
# again check if the no. of partitions are <=2
```

```
rdd_main.getNumPartitions()
```

```
# get the length of all IDs put together
```

```
recLenRdd = rddCoalesed.map(lambda rec: len(rec))
```

```
totalLength = recLenRdd.reduce(lambda a, b: a + b)
```

```
print("Total length of all strings:" + str(totalLength))
```

```
#Sort the IDs alphabetically and print the first 10
```

```
sortedRdd = rddCoalesed.map(lambda x: (x, 1)).sortByKey().keys()
```

```
sortedRdd.take(10)
```

```
# Store the sorted RDD in gs
```

```
sortedRdd.saveAsTextFile("gs://kamal_at_ie/sorted_rddids.txt")
```

```
#From rdd_main obtain the number of ids that share the first 2 chars
```

```
#For example if we have ids: aax , aat, aaron, bbt we would have aa,3 and bb,1
```

```
rddWordCount = rdd_main.map(lambda x: (x[:2], 1)).reduceByKey(lambda a, b: a + b)
```

```
# print the first 10
```

```
rddWordCount.take(10)
```